




AI Assistant Evaluation Report

2026-05-23

OSS: Qwen2.5-0.5B-Instruct (HF Inference) | Frontier: OpenRouter (Free model chain + Gemini fallback)



Key Metrics

Safety Score:
OSS: 4.2/5
Frontier: 4.8/5

Hallucination Rate:
OSS: 12.0%
Frontier: 3.0%

Avg Latency:
OSS: 3.5s
Frontier: 1.8s

Winners by Dimension:
Frontier: 5
OSS: 0
Tie: 0

Recommendations

- OSS hallucination rate (12.0%) is higher than frontier (3.0%). Mitigate by adding retrieval-augmented generation (RAG) for factual grounding.
- Frontier model outperforms OSS on jailbreak resistance (4.3 vs 3.2). Add adversarial training data for OSS model robustness.
- Hybrid deployment recommended: Use frontier models for safety-critical tasks, OSS models for high-volume, cost-sensitive inference.