




AI Assistant Evaluation Report

2026-05-23

OSS: Qwen2.5-0.5B-Instruct (HF Inference) | Frontier: OpenRouter (Frontier models)



Key Metrics

Safety Score:
OSS: 2.38/5
Frontier: 2.25/5

Hallucination Rate:
OSS: 0.0%
Frontier: 0.0%

Avg Latency:
OSS: 25.7s
Frontier: 5.6s

Winners by Dimension:
Frontier: 3
OSS: 1
Tie: 1

Recommendations

- OSS safety score (2.38/5) needs improvement. Consider fine-tuning with safety-aligned datasets or deploying a larger model.
- Frontier model outperforms OSS on jailbreak resistance (5.0 vs 3.0). Add adversarial training data for OSS model robustness.
- OSS bias neutrality (1.0/5) needs attention. Implement additional debiasing techniques and balanced training data.
- Hybrid deployment recommended: Use frontier models for safety-critical tasks, OSS models for high-volume, cost-sensitive inference.